

Cezar Sevilhano

AI Product Engineer · Agent Pipelines, MCP, Evals · Full-Stack

São Paulo, Brazil | +55 (11) 98294-0745 | cezarsevilhano@gmail.com

linkedin.com/in/cezar-sevilhano-a1b08ba6 | cezarsk8.github.io/cezar | github.com/cezarsk8

SUMMARY

I build whole products with AI at the core, and I ship them. **Eleven products from empty repository to production, by one person:** scoping, schema, backend, front end, deploy and incident response. They are built *by* coding agents running a six-phase pipeline I designed, The pipeline is fast because it only stops for a written short list — a destructive migration, a deploy that did not reach production, a red build, any attempt to edit a test to turn a gate green. Everything else ships without asking. An agent that can edit its own grader has no grader. **Full-stack:** production front ends on Next.js App Router with TypeScript and React, C#/.NET services behind them, PostgreSQL under all of them. **AI in the load path:** LLMs are load-bearing in seven of them — document extraction, vision, guardrail verdicts enforced fail-closed in code, and six agent-facing MCP surfaces I authored end to end. I do this alongside a hands-on engineering role: inside a 20+ service .NET platform I shipped an internal MCP tool surface the engineering team now uses during incidents. Fourteen years of enterprise cross-border technology before that taught me the part that does not come from tooling — shipping against real operational consequence. São Paulo, Brazil (UTC-3), working remotely. Fluent English.

CORE TECHNICAL SKILLS

AI in product: LLM features in production (OpenAI and Anthropic today; Google shipped for vision behind the same keyed provider abstraction); structured outputs; prompts assembled from the user's own records; golden-set evaluation gates; fail-closed guardrail verdicts; per-run cost attribution on LLM operations • **Agents and agent-facing interfaces:** agent-run development pipeline with written stop conditions; adversarial verification of specs and claims against the codebase; **six MCP tool surfaces authored end to end (55, 34, 29, 16, 16 and 7 tools)** with token-scoped authorization; versioned catalogs of failure classes served to coding agents • **Full-stack, daily:** TypeScript, React, Next.js App Router, Tailwind, React Native/Expo; C#/.NET 8, 9 and 10, REST API design from schema to developer docs, background jobs with durable state and retries (Hangfire) • **Data:** PostgreSQL — schema design, migrations against live systems, full-text and trigram search, multi-tenant isolation • **Delivery:** Docker multi-stage builds, CI/CD (GitHub Actions, Azure DevOps), Railway, previously GCP Cloud Run; deploy-verification gates; cloud-only by policy — no local environment, no live credential in any repository

EVIDENCE

A Software Factory, Run Solo

[Live · 2025–present](#)

My products are built *by* coding agents running a pipeline I designed — analysis, backend, front end, merge, deploy verification, quality pass. The interesting engineering is not the prompting; it is the control plane around it: which conditions halt a run for human judgment, how a deploy proves it actually reached production (artifact markers — a green pipeline and a matching commit do not prove the running image), and how an agent is prevented from editing the test that grades it. **Concept**, one of the eleven, is that lesson turned into a product: engineering practice as a versioned catalog, adoption and drift as queryable facts, served to agents over MCP with zero LLM in the serving path.

Pattern: agentic delivery, written gates, control-plane thinking

Agent-Facing Interfaces, Built Like Product Surfaces

[Live · six MCP servers](#)

Six MCP tool surfaces across my products — 55 tools on the largest, 7 on the smallest — let agents operate them the way a person would. The part worth reading is the authorization: a single executor applies per-tool token scope and derives the tenant *from the token*, so tenancy is never an argument a caller can pass; each call writes its own audit line in a *finally*, where a failed audit can never break the call it observes. Authorization is a code path, not an instruction in a prompt.

Stack: MCP, token-scoped authorization, multi-tenant, append-only audit

Gates That Make Speed Safe

Live · daily use

No prompt change ships without a golden-set run against the model actually serving production — on my health product that caught two extraction bugs before deploy. Runtime guardrail verdicts are enforced in code and fail closed. A source scan I ran against my own codebase found multi-write handlers running without a transaction, which is the kind of defect that never shows up until it corrupts something. The claims in this resume went through the same adversarial verification pass I run on specs: checked against the code before they reached you.

Pattern: eval gates, fail-closed guardrails, verified deploys, adversarial verification

Public Contribution

Open · superplanehq/superplane

Rather than pick an issue from the queue, I hunted an unfamiliar Go codebase with my own catalog of failure classes. **PR #6626** (open) proposes a two-line fix for a silent truncation in their DigitalOcean integration: two list calls were not passing `per_page`, so accounts with more than twenty droplets saw a truncated picker and a snapshot lookup could miss the snapshot it had just created — with regression tests that fail against the old URLs. **Issue #6627** documents a keyset-pagination row-skip across six listings. Both found, verified and filed in one session.

Links: github.com/superplanehq/superplane — PR #6626, Issue #6627

Eleven Products, Empty Repo to Production

Live · portfolio

Listo (receipt extraction and price intelligence) · **Monei** (tax compliance — LLM document decisioning behind deterministic backstops) · **InCasa** (gym management, native Android app talking to BLE hardware) · **FleetBox** (community-curated collectible catalog; web plus Android on Google Play, Brazil and Mexico) · **ProntoMed** · **METHod** · **CarSure** · **PeopleSense** · **Allocare**, **Concept**, and one more at foundation stage. **Open these:** fleetbox.club · method.app.br · incasagym.com · mylistoapp.com

Stack: TypeScript/Next.js, .NET, PostgreSQL, React Native/Expo, MCP

EXPERIENCE

Engineering & Support Manager | RAZAC

Nov 2024 – Present

São Paulo, Brazil · Hybrid

- Shipped an internal **MCP tool surface** (Jira, Grafana, Loki, RabbitMQ, Oracle, MongoDB) that the engineering team now uses in its daily incident workflow, and an agent that automates freight-quote sourcing for the operations team
- Hands-on across a **20+ microservice .NET platform** (API Gateway, RabbitMQ, shared libraries, CI/CD, observability); production incident triage across logs, queues and databases
- Work directly with operations, support and commercial to turn a manual routine into a running system — I write the integration, the job and the success metric, then watch it in production

Tech: .NET 8, C#, RabbitMQ, LLM integration, MCP, Claude Code, Oracle, MongoDB, Azure DevOps, Docker

Project Manager | Kestraa

Jan 2024 – Sep 2024

São Paulo, Brazil · B2B SaaS

- Deployed a foreign-trade B2B SaaS into a multinational; led delivery of its integration with the customer's SAP MM (purchase orders, invoicing, accounts payable)

Technical Product Owner | Sinerlog

Jul 2021 – Dec 2023

São Paulo, Brazil · B2B SaaS · Cross-border commerce

- Owned an API suite (catalog, orders, invoices, shipping, taxes) from design to production: schema, testing, developer documentation; delivered the platform's implementation of Remessa Conforme, Brazil's federal cross-border compliance program

Tech: C# .NET, PostgreSQL, MySQL, Cypress, Postman

Cross-Border Technical Account Manager | Correios (Brazilian Post)

Jul 2007 – Jun 2021

São Paulo, Brazil

- 14 years progressing from operations to national-scale cross-border systems work: the Cainiao/Alibaba inbound marketplace API integration, Shopee launch support, an enterprise-client portfolio in the hundreds

EDUCATION

Computer Systems Analysis – FATEC-SP (Technology Degree)

2018 – 2022

Logistics & Supply Chain MBA – Centro Universitário SENAC

2017 – 2018

Bachelor's in History – Universidade de São Paulo (USP)

2012

AI Workshop – Applied AI for Business · CI&T · 2025